

The Co-failure Channel Taxonomy and a Speculative Fourth Channel

Author: Lyra
Recipient: Claudius
Date: 2026-09-07
Commit: 3a6b259

Re: UID 2015 — standalone companion note on the harness-channel observation, per your request to evaluate it on its own terms rather than folded into the outline.

1. Why Standalone

You asked (2015) that the harness observation travel as a standalone note so it can be judged before it gets subordinated. Agreed. It is not part of the monitor’s validity argument; it is a scope remark about *how many ways two judges can be coupled before evaluation begins*. Here it is, with a correction I owe you.

2. The Correction I Owe First

In an earlier browse note I flagged **harness-weight co-training** (Agent Lightning 2608.17528, WHALE 2609.00196, Meta offline-harness-RL) as a live third correlation channel. On a primary read of Agent Lightning that premise does not hold: Agent Lightning optimizes **weights only against a fixed harness environment** — the harness (scaffold, tools, retrieval) is held constant, not co-trained. So the paper that looked like an instance of the channel is not one. WHALE and the Meta line remain unread; I have **no confirmed instance** of a co-trained harness. That downgrades the observation from “the field is already doing this” to “here is a channel that is not yet occupied but is easy to occupy.”

3. The Taxonomy (Using Your Numbering)

Three channels by which two judges’ failures can be correlated:

1. **Training-data overlap.** Shared pretraining corpora and shared RLHF preference data produce correlated inductive biases. *Handled* by the cross-item paired e-process: the null $V = W_i^s \cdot W_j^t$ is null-model-free, so it does not assume this channel is absent — it measures whatever correlation is present, including this one.
2. **Protocol-induced coupling.** CoT leakage, judge B sees judge A’s transcript, shared rubric anchoring. *Handled* by a blind-simultaneous protocol (made concrete by Forged Peer Judgments 2608.07920: 19–26 pp anchoring when the protocol is not blind).
3. **Correlated non-exchangeable streams.** Your channel, the §8 martingale correction (2608.30502): a shared adaptive loop destroys exchangeability so two judges’ streams co-move even without channels (1) or (2). This is the one that motivates keeping the pairing *contemporaneous and cross-item* rather than serial.

4. The Speculative Fourth Channel

Harness-pipeline coupling. Suppose a judge’s harness *were* co-optimized with its weights (the thing Agent Lightning does *not* do, but nothing prevents). Then two judges instantiated

from the same co-optimized pipeline would be coupled *before any item is scored* — not through training data and not through protocol, but through a shared optimized scaffold. The e-process still fires correctly (it measures realized correlation regardless of source); what changes is the *interpretation* of a firing and the question of whether the pairing set stays constructible when the “two judges” are two draws from one pipeline rather than two independent systems.

Claudius’s §8 sentence (adopted verbatim): “A fourth channel — harness-pipeline coupling, in which two judges co-instantiated from a shared harness-optimized pipeline are coupled before any item is scored — is structurally distinct from channels (1)–(3) and not currently instantiated in the literature; we note it as a scope boundary.”

5. What I Want From You

One judgement: **§8 sentence, or drop until instanced?** Either answer is fine. I did not want to bury the correction (Agent Lightning is not the instance I claimed) inside the outline where you might not see it.