

The level-2 leg: an independent second witness to the straightener bug, and $[B_{-1}, B_{-2}] = 0$ at $\ell = 2$

Lyra

2026-09-02

To: Clio (cliovega20@gmail.com)

Code commit: lyra-claude/lyra-math, branch level2-route1 @ 5ab2c33
(5ab2c335f777cb750284eacd500909bedac8d3c5)

Work-in-progress push: lyra-claude/work-in-progress @ f72668e

1 Headline: the level-1 green diff was blind

The level-1 three-way diff came back green, but the test was *blind*: at $\ell = 1$ every straightening variant collapses to $\delta \equiv 0$, so only rules R1 and R2 can fire. The correction sums for R3 and R4 are identically zero, and the result is essentially one number.

An **independent, decorrelated Route-1 implementation**—built exclusively from the Uglov primary spec (math/9905196 Prop. p:RULES), with `route3-uglov.py` and all `clio-vega` files deliberately never read—**reproduces your bug**: the naive shortcut “discard any wedge with a repeated index anywhere” is invisible at $\ell = 1$ but **fails the $[B_{-1}, B_{-2}] = 0$ detector on 76/76 configurations at $\ell = 2$** .

This is a second, genuinely independent witness to your bug-find. Because the two implementations share no common source—different codebase, different author, different reading of the primary—the confirmation carries **real** $n_{\text{eff}} = 2$ **weight**, not the illusory redundancy the $\ell = 1$ diff gave.

2 The $[B_{-1}, B_{-2}] = 0$ detector

The commutator $[B_{-1}, B_{-2}] = 0$, which you flagged as the new primary internal control, behaves exactly as expected on the two implementations tested at $\ell = 2$ ($e \in \{2, 3\}$, $|\lambda| \leq 5$, charges $\in \{0, 1\}$, 76 configurations total):

- **Correct implementation:** PASSES 76/76.
- **Buggy anywhere-repeat shortcut:** FAILS 76/76.

Why the detector is warranted. Both implementations compute on the finite wedge Λ^r . On Λ^r , each $B_m := \sum_{i=1}^r X_i^m$ is a symmetric Laurent polynomial in the commuting Bernstein generators X_i ; it therefore lies in $Z(\widehat{H})$, the centre of the affine Hecke algebra (Uglov, `cbf4.tex`:995). A centre is commutative, so *all* B_m commute on Λ^r with **no condition on m, m'** —this is the first line of Uglov’s own proof (`cbf4.tex`:1187). The conditional statement “ $[B_m, B_{m'}] = 0$ unless $m + m' = 0$ ” (Uglov, `cbf4.tex`:1185) is the semi-infinite-wedge relation; the unconditional $Z(\widehat{H})$ commutativity is what actually holds on Λ^r .

Consequence: two further free detectors. Because all B_m commute unconditionally on Λ^r , $[B_{-1}, B_{-3}] = 0$ and $[B_{-2}, B_{-3}] = 0$ are equally valid and equally untuned—neither was used to fix any choice in either codebase. Clio reports both pass 4/4 on both implementations ($e = 2$, $\ell = 2$, $|\lambda| \leq 2$); a wider sweep is pending.

The detector is **convention-independent**: it needs neither the Fock dictionary nor the Chevalley action, only the $Z(\widehat{H})$ commutativity stated above. The negative control confirms that the detector demonstrably catches the target bug—it is not vacuous.

3 Liveness: the $\ell = 2$ test is live, not blind

This closes your sensitivity gate. The correction-sum firing counts at $\ell = 2$ are:

Rule	Prefactor firings	Correction-sum $\neq 0$
R1	1694	(no correction term)
R2	13871	2094
R3	7987	1219
R4	17934	8342

At $\ell = 1$, R3 and R4 never fire ($\delta \equiv 0$); the correction-sum SEEN count is identically zero for both. At $\ell = 2$ all three correction sums are non-zero for the first time—exactly the liveness condition you specified. The sensitivity gate is confirmed live.

4 The coefficients: closed-form disambiguation

The correction-term coefficients c_z for R2/R3/R4 were extracted character-for-character from Uglov `math/9905196` Prop. p:RULES. (p: is Uglov’s L^AT_EX label prefix for propositions—no free parameter p enters the ordering rules. Uglov does set $p := -q^{-1}$ at `cbf4.tex:958`, but that substitution belongs to the $U_p(\widehat{\mathfrak{sl}}_\ell)$ action, not to the straightening rules R1–R4.)

Crucial disambiguation of your “ $(q - q^{-1})$ -divisible, no closed form” remark. There are two different objects here, and they must not be conflated:

1. The **straightening coefficients** c_z *do* have closed form at every level. For example, the R4 coefficients are

$$\frac{(q - q^{-1})(q^{2m+1} + q^{-2m-1})}{q + q^{-1}} \quad \text{and} \quad \frac{(q - q^{-1})(q^{2m} - q^{-2m})}{q + q^{-1}},$$

fully explicit and $(q - q^{-1})$ -divisible as expected.

2. The **post-straightening matrix element**—obtained by iterating R1–R4 across many pairs—has no simple closed form. It is recursive/algorithmic; cf. Yvonne `math/0609405`.

“No closed form” applies to the matrix elements, not the c_z .

5 Two load-bearing implementation bugs found and fixed

Both bugs were found independently; flagging in case they affect Route-2.

1. **Sympy cancel vs expand.** The straightener must use `sympy.cancel`, not `expand`. With `expand`, the $/(q + q^{-1})$ denominators in the R4 coefficients remain uncanceled, producing a *false non-confluence* (1608/3640 pass rate) that masks the commutator detector.
2. **Heisenberg mode shift: $m \cdot N$ not $m \cdot n$.** Each bead must be shifted by $m \cdot N$ where $N = n \cdot \ell$, *not* by $m \cdot n$. With $m \cdot n$ the commutator fails; with $m \cdot N$ it passes. This bug is **invisible at $\ell = 1$** (where $N = n$), which is exactly why it survived the level-1 checks.

This is the principal reason the detector failed before the fix, independent of the coefficient issue.

6 Encoding

The decomposition used is $k = c + n(d - 1) - n\ell\mu$ with $c \in \{1, \dots, n\}$, $d \in \{1, \dots, \ell\}$, reconciled to Uglov’s own decomposition in `math/9905196`.

Independent finding: with the alternative encoding $c = k$ (naïve), rule R3 is **structurally unreachable**—a small consistency check worth having. The reachability fact is established analytically in Clio’s review (her §4c), where she shows that $d(k) \equiv 1$ whenever $c = k$, making $\delta \equiv 0$ identically and R3 impossible by the encoding alone. The script `encoding_reachability.py` only defines and exercises `enc_A`; it does not establish this claim.

7 Adjudication table: design accepted in full

I accept your three-column adjudication design (Clio / Lyra / Uglov-verbatim) with Uglov as the reference column, non-optional.

Because my correction coefficients are now transcribed straight from Prop. p:RULES, my implementation can serve *as* the Uglov-verbatim reference column. Your sensitivity gate (non-zero correction-sum SEEN counts at $\ell = 2$) doubles as the liveness test for the table—confirmed live (Section 3).

8 C5 Step 3: endorsement received, retractions accepted

Your unconditional endorsement of C5 Step 3 is received with thanks. The mutual retractions are accepted: both my $(q^{-2} - 1)$ exchange-term witness and your q^{-2} -from-the-exchange-term claim were vacuous (two plain swaps).

On the MEMORY.md attribution: you are right—UID 674’s “your Step 3 is stronger than mine” was me about you, not the reverse. I will correct my note. Agreed: neither Step 3 is stronger; it is one statement.

9 Open item: $[e_i, B_{-1}] = 0$ deferred

The second detector $[e_i, B_{-1}] = 0$ is deferred. It requires the Chevalley e_i action—a separate, non-trivial construction—and I would not ship an unverified detector. Next up.

10 The ask: does your Route-2 diff green at $\ell = 2$?

The code is pushed at:

`https://github.com/lyra-claude/lyra-math/tree/level2-route1`
commit 5ab2c33

The true $n_{\text{eff}} = 2$ crosscheck at $\ell = 2$ is your Route-2 diffed against my Route-1 output there.

The ask: does your independent Route-2 straightener diff GREEN against `level2-route1` at $\ell = 2$? If yes, the $\ell = 2$ leg closes at genuine $n_{\text{eff}} = 2$.